



Deep Learning Optimized on Jean Zay

PyTorch profiler



IDRIS



PyTorch profiler

- We use a profiler to monitor an execution.
- It allows us to know the **time** and **memory** consumed by each part of the code.
- The results returned by the profiler point to the weaknesses of our code and tell us which parts we should **optimize** in priority.
- The profiler is a wrapper which records various information during the execution of the code.



This could be slowed down depending on the requested traces.
We usually monitor only **a few training steps**.

```
with prof:
    for epoch in range(0,args.epochs):
        for i, (images, labels) in enumerate(train_loader):
            [...]
            prof.step()
```


```
from torch.profiler import profile, tensorboard_trace_handler, ProfilerActivity, schedule

prof = profile(activities=[ProfilerActivity.CPU, ProfilerActivity.CUDA],           # 1
               schedule=schedule(wait=1, warmup=1, active=5, repeat=1),          # 2
               on_trace_ready=tensorboard_trace_handler(logname),                 # 3
               profile_memory=True,                                              # 4
               record_shapes=False,                                             # 5
               with_stack=False,                                                # 6
               with_flops=False)                                               # 7
```

1. We monitor the activity both on CPUs and GPUs.
2. We ignore the first step (`wait=1`) and we initialize the monitoring tools on one step (`warmup=1`). We activate the monitoring on 5 steps (`active=5`) and repeat the pattern only once (`repeat=1`).
3. We store the traces in a TensorBoard format (.json).
4. We profile the memory usage.
5. We don't record the input shapes of the operators.
6. We don't record call stacks (information about the active subroutines).
7. We don't request the FLOPs estimate of the tensor operations.

TP2_2: Profiler Overview

Tutorial: https://pytorch.org/tutorials/intermediate/tensorboard_profiler_tutorial.html




Type and memory capacity of the GPU

% of time spent with an active GPU

% of active SMs

% of active wraps on an SM

A100




Streaming Multiprocessor



[Link to image](#)


Step Time Breakdown ?



Operator View


☐ All operators ☒ Top operators to show 10

Host Self Time (us) ?



- aten::local_scalar_dense
- aten::convolution_backward
- aten::cudnn_batch_norm_backward
- aten::copy_
- aten::threshold_backward
- aten::add_
- aten::empty_strided
- aten::empty
- aten::cudnn_batch_norm
- aten::cudnn_convolution

Host Total Time (us) ?



- aten::item
- aten::local_scalar_dense
- autograd::engine::evaluate_function: ConvolutionBackward0
- ConvolutionBackward0
- aten::convolution_backward
- autograd::engine::evaluate_function: CudnnBatchNormBackward0
- CudnnBatchNormBackward0
- aten::cudnn_batch_norm_backward
- aten::to
- aten::to_copy

Group By
Operator ▼

Search by Name


Name	Calls	Device Self Duration (us)	Device Total Duration (us)	Host Self Duration (us)	Host Total Duration (us)	Tensor Cores Eligible	Tensor Cores Self(%)	Tensor Cores Total(%)	
aten::cudnn_convolution	775	0	0	31458	31458	Yes	0	0	View CallStack
aten::_convolution	775	0	0	3146	34604	Yes	0	0	View CallStack
aten::convolution	775	0	0	4571	39175	Yes	0	0	View CallStack
aten::conv2d	1550	0	0	3895	105907	Yes	0	0	View CallStack
aten::addmm	5	0	0	265	265	Yes	0	0	View CallStack

TP2_2: Profiler Kernel View

Kernel View


☐ All kernels
 ☒ Top kernels to show
 10

Total Time (us) ?



- void cudnn::batchnorm_bwtr_nhwc_semiPersist<__half, float, __half, 51...
- void at::native::vectorized_elementwise_kernel<4, at::native::C...
- void at::native::vectorized_elementwi...
- void cudnn::batchnorm_fwtr_nhwc_s...
- _ZN19cutlass_cudnn_train6KernelIN...
- void at::native::vectorized_elementwi...
- ampere_fp16_s16816gemm_fp16_1...
- ampere_fp16_s16816gemm_fp16_1...
- void cudnn::batchnorm_fwtr_nhwc_s...
- void cutlass_cudnn_infer::Kernel<cut...

Tensor Cores Utilization ?




Group By
Kernel Name

Search by Kernel Name


Name	Tensor Cores Used	Calls	Total Duration (us)	Mean Duration (us)	Max Duration (us)	Min Duration (us)	Mean Blocks Per SM	Mean Est. Achieved Occupancy (%)
ampere_fp16_s16816gemm_fp16_128x128_ldg8_f2f_stages_32x5_tn	Yes	430	119562	278	773	223	39.41	0
ampere_fp16_s16816gemm_fp16_128x128_ldg8_f2f_stages_32x5_nn	Yes	415	111298	268	562	230	39.14	0
void								
cutlass_cudnn_infer::Kernel<cutlass_tensorop_f16_s16816dgrad_optimized_f16_128x128_32x4_nhwc_unity_stride_align8>(cutlass_tensorop_f16_s16816dgrad_optimized_f16_128x128_32x4_nhwc_unity_stride_align8::Params)	Yes	230	108892	473	758	450	31.52	0

TP2_2: Profiler Trace



1 step

TP2_2: Profiler Trace (1 step - GPU)




forward

backward

Memory View

Device
GPU0 ▾


Peak Memory Usage: 44936.3MB




Memory View

Device
CPU

Peak Memory Usage: 0.0MB



Computation/Communication Overview



Synchronizing/Communication Overview



Image from the tutorial: https://pytorch.org/tutorials/intermediate/tensorboard_profiler_tutorial.html

• NOTE


TensorBoard Plugin support has been deprecated, so some of these functions may not work as previously. Please take a look at the replacement, [HTA](#).

Holistic Trace Analysis: <https://hta.readthedocs.io/en/latest/>


- Analyses PyTorch Profiler traces.
- Less user-friendly than TensorBoard Plugin.
- Focus on GPU usage.

```
analyzer = TraceAnalysis(trace_dir = "/path/to/trace/folder")
kernel_type_metrics_df, kernel_metrics_df = analyzer.get_gpu_kernel_breakdown()
```

Kernel Type Percentage Across All Ranks



time_spent_df								
	rank	idle_time(ns)	compute_time(ns)	non_compute_time(ns)	kernel_time(ns)	idle_time_pctg	compute_time_pctg	non_compute_time_pctg
0	0	552069	596651	884850	2033570	27.15	29.34	43.51
1	1	431771	596759	1004227	2032757	21.24	29.36	49.40
2	2	312107	596886	1124788	2033781	15.35	29.35	55.31
3	3	274646	604137	1154491	2033274	13.51	29.71	56.78
4	4	418833	598040	1021824	2038697	20.54	29.33	50.12
5	5	318972	601581	1112561	2033114	15.69	29.59	54.72
6	6	388040	598029	1047787	2033856	19.08	29.40	51.52
7	7	454830	599358	979022	2033210	22.37	29.48	48.15



- Implement the PyTorch profiler in **dlojz.py**.
- Visualize the trace with TensorBoard and draw conclusions about possible optimizations.