



Deep Learning Optimized on Jean Zay

Distribution – Data parallelism




IDRIS



Distributed training


General knowledge about parallel computing ◀

Data parallelism to distribute your training ◀



Sequential execution

- Only one process executes the program.
- The variables defined in the program are stored in the memory allocated to the process.
- One process executes the code on one physical compute unit (CPU core or GPU).




Parallel execution with distributed memory

- Several processes execute the code at the same time (multiprocessing).
- The variables defined in the program are private, they are stored in the local memory allocated to each process.
- It is possible that the processes execute separate parts of the code.


Parallel execution with distributed memory

- To share information, the processes can send each other **messages** through the interconnection network.
- These **communications** are managed by libraries such as **MPI** or **NCCL**.



In a parallel code based on the MPI or NCCL library,

- The set of processes exists in a common parallel **environment** initialized at the beginning of the execution.
- From the initialization to the destruction of this parallel environment, **all the processes** read and execute the program.
- During the initialization of the environment, a **communicator WORLD** is created to allow the set of processes to communicate with each other.
- The communicator has a **size** (the number of processes).
- Within a communicator, each process can be identified by its **rank**.



Concrete examples based on the Horovod library

- The **Horovod** library is designed to ease the implementation of Deep Learning distributed training

Communication libraries	
•MPI (on CPU)	✓
•NCCL (on GPU)	✓

Deep learning frameworks	
•TensorFlow	✓
•Keras	✓
•PyTorch	✓
•Apache MXNet	✓

Example 1: Each process reads and executes the code lines

- Parallelized code using Horovod:

```
import horovod.torch as hvd  
  
hvd.init()  
size = hvd.size()  
print(f'The communicator size is {size}')
```

- Execution on 4 processes (#SBATCH --ntasks=4):

```
$ srun --ntasks=4 <...> python script.py  
The communicator size is 4  
The communicator size is 4  
The communicator size is 4  
The communicator size is 4
```

Example 2: We identify the processes thanks to their ranks

- Parallelized code using Horovod:

```
import horovod.torch as hvd

hvd.init()
size = hvd.size()
rank = hvd.rank()
print(f'I am proc {rank} among {size}')
```

- Execution on 4 processes (#SBATCH --ntasks=4):

```
$ srun --ntasks=4 <...> python script.py
I am rank 1 among 4
I am rank 3 among 4
I am rank 2 among 4
I am rank 0 among 4
```

Example 3: The processes can be assigned different tasks according to their ranks

- Parallelized code using Horovod:

```
import horovod.torch as hvd

hvd.init()
size = hvd.size()
rank = hvd.rank()
print(f'I am proc {rank}, my rank is {"even" if rank%2==0 else "odd"}')
```

- Execution on 4 processes (#SBATCH --ntasks=4):

```
$ srun --ntasks=4 <...> python script.py
I am proc 2, my rank is even
I am proc 0, my rank is even
I am proc 1, my rank is odd
I am proc 3, my rank is odd
```

Example 3bis: The processes can be assigned different tasks according to their ranks

- Parallelized code using Horovod:

```
import horovod.torch as hvd

hvd.init()
size = hvd.size()
rank = hvd.rank()
if rank%2==0:    print(f'I am proc {rank}, my rank is even')
else:           print(f'I am proc {rank}, my rank is odd')
```

- Execution on 4 processes (#SBATCH --ntasks=4):

```
$ srun --ntasks=4 <...> python script.py
I am proc 2, my rank is even
I am proc 3, my rank is odd
I am proc 0, my rank is even
I am proc 1, my rank is odd
```

Example 4: Parallelization of a compute loop

- Initial state of the memory
- Program
- Final state of the memory

↓Process 0↓

```
N = 4  
a = [0, 1, 2, 3]  
b = [4, 5, 6, 7]  
c = [0, 0, 0, 0]  
sum = 0
```

```
for i in range(N) :  
    c[i] = a[i] + b[i]  
    sum += c[i]
```

```
N = 4  
a = [0, 1, 2, 3]  
b = [4, 5, 6, 7]  
c = [4, 6, 8, 10]  
sum = 28
```

Example 4: Parallelization of a compute loop

- Initial state of the memory

- Program

- Final state of the memory

↓Process 0↓

```
N = 4
a = [0, 1, 2, 3]
b = [4, 5, 6, 7]
c = [0, 0, 0, 0]
sum = 0
```

↓Process 1↓

```
N = 4
a = [0, 1, 2, 3]
b = [4, 5, 6, 7]
c = [0, 0, 0, 0]
sum = 0
```

```
size = hvd.size()
rank = hvd.rank()

istart = rank * N / size
iend = (rank+1) * N / size

for i in range(istart,iend) :
    c[i] = a[i] + b[i]
    sum += c[i]
```

```
N = 4
a = [0, 1, 2, 3]
b = [4, 5, 6, 7]
c = [4, 6, 0, 0]
sum = 10
```

```
N = 4
a = [0, 1, 2, 3]
b = [4, 5, 6, 7]
c = [0, 0, 8, 10]
sum = 18
```

Example 4: Parallelization of a compute loop

- Initial state of the memory

- Program

- Final state of the memory

↓Process 0↓

```
N = 4
a = [0, 1, 2, 3]
b = [4, 5, 6, 7]
c = [0, 0, 0, 0]
sum = 0
```


↓Process 1↓

```
N = 4
a = [0, 1, 2, 3]
b = [4, 5, 6, 7]
c = [0, 0, 0, 0]
sum = 0
```

```
size = hvd.size()
rank = hvd.rank()

istart = rank * N / size
iend = (rank+1) * N / size
```



```
for i in range(istart,iend) :
    c[i] = a[i] + b[i]
    sum += c[i]
sum = hvd.allreduce(sum,op=hvd.Sum)
```



```
N = 4
a = [0, 1, 2, 3]
b = [4, 5, 6, 7]
c = [4, 6, 0, 0]
sum = 28
```

```
N = 4
a = [0, 1, 2, 3]
b = [4, 5, 6, 7]
c = [0, 0, 8, 10]
sum = 28
```

- Inter-process communication `AllReduce()`
 - Collective communication
 - Synchronization barrier









AllReduce
($\Sigma, \Pi, \min, \max$)




costly communication


- NCCL communications



Sequential execution




Parallel execution



- Global execution time:


Scalability study:

- **Strong scaling**
(problem size is constant)
- **Weak scaling**
(problem size is proportional to the number of processes)




(1) impact of the sequential parts of the code
(2) cost of the inter-process communications

- Bandwidths of the interconnection networks on Jean Zay:




Node 4 × V100 16GB

Flaubert benchmark




Strong scaling

- Bandwidths of the interconnection networks on Jean Zay:



Node 8 × A100 80GB

Flaubert benchmark



Strong scaling

Distributed training


General knowledge about parallel computing ◀

Data parallelism to distribute your training ◀

Data parallelism

- Training time speedup
- Model small enough to be contained on one GPU in memory
 - Causes large batches (consequences on the training quality)


Distribution: Data parallelism



Implementation of the data parallelism

- PyTorch → **DistributedDataParallel** (integrated solution)
- TensorFlow → `MultiWorkerMirroredStrategy` (integrated solution)
- Horovod (external library)

Distribution: Data parallelism



DistributedDataParallel


Distribution: Data parallelism

Execution of the parallel code → Slurm environment

- Distribution example on : 4 nodes
3 GPUs per node

```
## Slurm script
#SBATCH --nodes=4          # nb nodes
#SBATCH --ntasks=12        # nb proc
#SBATCH --ntasks-per-node=3 # nb proc / node
#SBATCH --gres=gpu:3       # nb GPUs / node
srun python script.py
```

Each GPU must be bound with one process.



- **DistributedDataParallel** (training distribution)

```
import torch.distributed
from torch.nn.parallel import DistributedDataParallel
```

```
# initialize the parallel environment
```


```
torch.distributed.init_process_group(backend='nccl',
                                     init_method='env://',
                                     world_size=world_size,
                                     rank=rank)
```

```
# bind one GPU per process
```

```
torch.cuda.set_device(local_rank)
```

```
# duplicate the model
```

```
ddp_model = DistributedDataParallel(model, device_ids=[local_rank])
```



Distribution: Data parallelism

- **DistributedDataParallel** (training distribution)
 - *idr_torch.py* script from IDRIS

```
# idr_torch.py
import os
import hostlist

# get SLURM variables
size = int(os.environ['SLURM_NTASKS'])
rank = int(os.environ['SLURM_PROCID'])
local_rank = int(os.environ['SLURM_LOCALID'])
cpus_per_task = int(os.environ['SLURM_CPUS_PER_TASK'])

# get node list from slurm
hostnames = hostlist.expand_hostlist(os.environ['SLURM_JOB_NODELIST'])

# get IDs of reserved GPU
gpu_ids = os.environ['SLURM_STEP_GPUS'].split(",")

# define MASTER_ADD & MASTER_PORT
os.environ['MASTER_ADDR'] = hostnames[0]
os.environ['MASTER_PORT'] = str(12345 + int(min(gpu_ids)))
```

- **DistributedDataParallel** (training distribution)

```
import idr_torch
import torch.distributed
from torch.nn.parallel import DistributedDataParallel

# initialize the parallel environment
torch.distributed.init_process_group(backend='nccl',
                                    init_method='env://',
                                    world_size=idr_torch.size,
                                    rank=idr_torch.rank)

# bind one GPU per process
torch.cuda.set_device(idr_torch.local_rank)

# duplicate the model
ddp_model = DistributedDataParallel(model, device_ids=[idr_torch.local_rank])
```

MASTER_PORT

MASTER_ADDR

SLURM_NTASKS

SLURM_PROCID

SLURM_LOCALID

Slurm

Distribution: Data parallelism

- What about torchrun? → Possible but cumbersome.


```
[...]  
#SBATCH --ntasks-per-node=1  
[...]  
GPUS_PER_NODE=8  
MASTER_ADDR=$(scontrol show hostnames $SLURM_JOB_NODELIST | head -n 1)  
MASTER_PORT=15000  
CMD="train.py --arg1 1 --arg2 2"  
export LAUNCHER="torchrun --nproc_per_node $GPUS_PER_NODE \  
                --nnodes $SLURM_NNODES \  
                --rdzv_backend c10d \  
                --rdzv_endpoint $MASTER_ADDR:$MASTER_PORT"  
srun bash -c "$LAUNCHER --node_rank $SLURM_PROCID $CMD"
```

slurm.sh

```
[...]  
WORLD_RANK = int(os.environ['RANK'])  
LOCAL_RANK = int(os.environ['LOCAL_RANK'])  
WORLD_SIZE = int(os.environ['WORLD_SIZE'])
```

train.py

Distribution: Data parallelism



- **DistributedSampler** (distributing the input data)

```
import idr_torch
import torch.distributed
from torch.nn.parallel import DistributedDataParallel
from torch.utils.data.distributed import DistributedSampler
```

```
# initialize the parallel environment
[...]
```

```
# bind one GPU per process
[...]
```

```
# duplicate the model
[...]
```

```
# distribute the input data
```

```
data_sampler = DistributedSampler(dataset, shuffle=True,
                                  num_replicas=idr_torch.size, rank=idr_torch.rank)
```

SLURM_NTASKS

Slurm

SLURM_PROCID

The shuffling step is assigned to the Sampler.

- **DistributedSampler** (distributing the input data)
- Parallel execution on 4 processes using the DistributedSampler:

```
{ dataset = [0, 1, ..., 99]  
  batch_size_per_gpu = 5  
  ntasks = 4  
  batch_size = 20
```

```
$ srun --ntasks=4 <...> python script.py  
Rank 0: Batch 0 = tensor([ 0,  4,  8, 12, 16])  
Rank 1: Batch 0 = tensor([ 1,  5,  9, 13, 17])  
Rank 2: Batch 0 = tensor([ 2,  6, 10, 14, 18])  
Rank 3: Batch 0 = tensor([ 3,  7, 11, 15, 19])
```

- **DistributedSampler** (distributing the input data) + **shuffling**
 - The index shuffling is performed by each GPU from a common seed.
- Parallel execution on 4 processes using the DistributedSampler:

```
{ dataset = [0, 1, ..., 99]
  batch_size_per_gpu = 5
  ntasks = 4
  batch_size = 20
```

```
$ srun --ntasks=4 <...> python script.py
Rank 0: Batch 0 = tensor([46, 36, 80, 17, 23])
Rank 1: Batch 0 = tensor([16, 64, 97, 12, 59])
Rank 2: Batch 0 = tensor([91, 18, 49, 24, 4])
Rank 3: Batch 0 = tensor([33, 73, 37, 81, 63])
```

- **DistributedSampler** (distributing the input data) + **shuffling**
 - The index shuffling is performed by each GPU from a common seed.



```
for epoch in range(1,30):  
    for i, batch in enumerate(dataloader):  
        ...
```

```
>>> Epoch 1  
Rank 0: Batch 0 = tensor([46, 36, 80, 17, 23])  
Rank 1: Batch 0 = tensor([16, 64, 97, 12, 59])  
Rank 2: Batch 0 = tensor([91, 18, 49, 24, 4])  
Rank 3: Batch 0 = tensor([33, 73, 37, 81, 63])  
>>> Epoch 2  
Rank 0: Batch 0 = tensor([46, 36, 80, 17, 23])  
Rank 1: Batch 0 = tensor([16, 64, 97, 12, 59])  
Rank 2: Batch 0 = tensor([91, 18, 49, 24, 4])  
Rank 3: Batch 0 = tensor([33, 73, 37, 81, 63])
```

Distribution: Data parallelism

- **DistributedSampler** (distributing the input data) + **shuffling**
 - The index shuffling is performed by each GPU from a common seed.



```
for epoch in range(1,30):  
    data_sampler.set_epoch(epoch)  
    for i, batch in enumerate(dataloader):  
        ...
```

```
>>> Epoch 1  
Rank 0: Batch 0 = tensor([46, 36, 80, 17, 23])  
Rank 1: Batch 0 = tensor([16, 64, 97, 12, 59])  
Rank 2: Batch 0 = tensor([91, 18, 49, 24, 4])  
Rank 3: Batch 0 = tensor([33, 73, 37, 81, 63])  
>>> Epoch 2  
Rank 0: Batch 0 = tensor([49, 91, 8, 76, 48])  
Rank 1: Batch 0 = tensor([98, 50, 21, 15, 22])  
Rank 2: Batch 0 = tensor([ 2, 11, 71, 92, 75])  
Rank 3: Batch 0 = tensor([82, 9, 74, 39, 53])
```

- Custom Sampler (inspired by DistributedSampler)

```
class MyCustomDistributedSampler(Sampler):
```

```
    def __init__(self, dataset, world_size, rank):
        self.data_len = len(dataset)
        self.world_size = world_size
        self.rank = rank
```


```
    def __len__(self):
        return self.data_len
```

```
    def __iter__(self):
        indices = list(range(self.data_len))
        # shuffle or not shuffle
        indices = indices[self.rank:self.data_len:self.world_size]
        return iter(indices)
```


```
sampler = MyCustomDistributedSampler(dataset, idr_torch.size, idr_torch.rank)
```

```
$ srun --ntasks=4 <...> script.py
Rank 0: Batch 0 = tensor([ 0,  4,  8, 12, 16])
Rank 1: Batch 0 = tensor([ 1,  5,  9, 13, 17])
Rank 2: Batch 0 = tensor([ 2,  6, 10, 14, 18])
Rank 3: Batch 0 = tensor([ 3,  7, 11, 15, 19])
```





- Go into the directory `tp_pi/`
- Follow instructions in the notebook **DL0-JZ_Compute_pi.ipynb**
- Parallelize the code `compute_pi.py` using `torch.distributed`
- Compute π on 4 GPUs



- Follow instructions in the notebook **DL0-JZ_Jour2.ipynb**
- Implement data parallelism in the script `dlojz.py`
- Measure the gain in time when using 4 GPUs