




# Optimised Deep Learning on Jean Zay (DLO-JZ)


---

## Torch & Compilation



INSTITUT DU  
DÉVELOPPEMENT ET DES  
RESSOURCES EN  
INFORMATIQUE  
SCIENTIFIQUE





```

1 #include <stdio.h>
2
3 int main (void) {
4     printf ("Hello, World!\n");
5
6     return 0;
7 }

```

```

hello.py
hello.py > ...
1 msg = "Hello World"
2 print(msg)

```

Compiled

C, C++, Go, Fortran, Pascal

Language

"Compiling"

Machine Code

Ready to Run!

Interpreted

Python, PHP, Ruby, JavaScript


Language




Ready to Run!


"Interpreting"


Virtual Machine

Machine Code









**Internal Pytorch JIT:**

**TorchScript**

**Torch Compile**



```
import torch
class MyCell(torch.nn.Module):
    def __init__(self):
        super(MyCell, self).__init__()
        self.linear = torch.nn.Linear(4, 4)

    def forward(self, x, h):
        new_h = torch.tanh(self.linear(x) + h)
        return new_h, new_h

x, h = torch.rand(3, 4), torch.rand(3, 4)
my_cell = MyCell()
my_cell(x, h)
```

## Tracing

```
traced_cell = torch.jit.trace(my_cell, (x, h))
traced_cell(x, h)
print(traced_cell.code)

def forward(self,
    x: Tensor,
    h: Tensor) -> Tuple[Tensor, Tensor]:
    linear = self.linear
    _0 = torch.tanh(torch.add((linear).forward(x, ), h))
    return (_0, _0)
```

```

class MyDecisionGate(torch.nn.Module):
    def forward(self, x):
        if x.sum() > 0:
            return x
        else:
            return -x

class MyCell(torch.nn.Module):
    def __init__(self, dg):
        super(MyCell, self).__init__()
        self.dg = dg
        self.linear = torch.nn.Linear(4, 4)

    def forward(self, x, h):
        new_h = torch.tanh(self.dg(self.linear(x)) + h)
        return new_h, new_h

my_cell = MyCell(MyDecisionGate())
traced_cell = torch.jit.trace(my_cell, (x, h))

```

**TracerWarning: Converting a tensor to a Python boolean might cause the trace to be incorrect. We can't record the data flow of Python values, so this value will be treated as a constant in the future. This means that the trace might not generalize to other inputs!**

## Scripting

```

scripted_gate = torch.jit.script(MyDecisionGate())

my_cell = MyCell(scripted_gate)
scripted_cell = torch.jit.script(my_cell)

```

```
print(f"Non-traced average time: {non_traced_time:.6f} seconds")
print(f"Traced average time: {traced_time:.6f} seconds")
print(f"Scripted average time: {traced_time:.6f} seconds")
```

```
Non-traced average time: 0.000030 seconds
Traced average time: 0.000038 seconds
Scripted average time: 0.000038 seconds
```

|             | Latency on CPU (ms) | Latency on GPU(ms) |
|-------------|---------------------|--------------------|
| PyTorch     | 86.23               | 16.49              |
| TorchScript | 81.57               | 10.54              |

PyTorch vs TorchScript for BERT

|             | Latency on CPU (ms) | Latency on GPU(ms) |
|-------------|---------------------|--------------------|
| PyTorch     | 25.96               | 4.02               |
| TorchScript | 23.01               | 2.41               |

PyTorch vs TorchScript for ResNet


### Save in Python


```
traced_cell.save('wrapped_cell.pt')
```

### Load in C++

```
#include <torch/script.h>
...
torch::jit::script::Module module;
...
module = torch::jit::load(model_path);
...
```

forward





```
import torch
def fn(a, b, c, d):
    x = a + b + c + d
    return x.cos().cos()
a, b, c, d = [torch.randn(2, 4, requires_grad=True) for _ in range(4)]
ref = fn(a, b, c, d)
loss = ref.sum()
loss.backward()
```

```
from functorch.compile import aot_function

def compiler_fn(fx_module: torch.fx.GraphModule, _):
    print(fx_module.code)
    return fx_module


aot_print_fn = aot_function(fn, fw_compiler=compiler_fn, bw_compiler=compiler_fn)




# To trigger the compilation and see the result:
res = aot_print_fn(a, b, c, d)

def forward(self, primals_1, primals_2, primals_3, primals_4):
    add = torch.ops.aten.add.Tensor(primals_1, primals_2); primals_1 = primals_2 = None
    add_1 = torch.ops.aten.add.Tensor(add, primals_3); add = primals_3 = None
    add_2 = torch.ops.aten.add.Tensor(add_1, primals_4); add_1 = primals_4 = None
    cos = torch.ops.aten.cos.default(add_2)
    cos_1 = torch.ops.aten.cos.default(cos)
    return [cos_1, add_2, cos]
```

```
torch.compile(model=None, *, fullgraph=False, dynamic=None, backend='inductor', mode=None,
options=None, disable=False) \[SOURCE\]
```

## TorchDynamo + AOT AutoGrad + TorchInductor + PrimTorch







| Compilation mode | Initial step time [s] | Initial step memory [MiB] | Training time / iter [ms] | Inference time / iter [ms] |
|------------------|-----------------------|---------------------------|---------------------------|----------------------------|
| N/A (eager)      | 1                     | 3675                      | 57                        | 18                         |
| default          | 29                    | 3277                      | 34                        | 30                         |
| reduce-overhead  | 29                    | 5736                      | 34                        | 28                         |
| max-autotune     | 35                    | 5736                      | 32                        | 30                         |



```
import numpy as np

def kmeans(X, means):
    return np.argmin(np.linalg.norm(X - means[:, None], axis=2), axis=0)

npts = 10_000_000
X = np.repeat([[5, 5], [10, 10]], [npts, npts], axis=0)
X = X + np.random.randn(*X.shape) # 2 distinct "blobs"
means = np.array([[5, 5], [10, 10]])
np_pred = kmeans(X, means)
```

### Torch Compile Wrapping

```
import torch

compiled_fn = torch.compile(kmeans)
compiled_pred = compiled_fn(X, means)
assert np.allclose(np_pred, compiled_pred)
```

### Torch Compile Performance

```
print(f"Non-compiled average time: {non_compiled_time:.6f} seconds")
print(f"Compiled average time: {compiled_time:.6f} seconds")
```

```
Non-compiled average time: 2.521456 seconds
Compiled average time: 0.233836 seconds
```



**RAPIDS**

Accelerated with



**CuPy**





# Flax Linen

*Neural networks with JAX*

Flax Linen delivers an **end-to-end and flexible user experience for researchers who use JAX with neural networks**. Flax exposes the full power of [JAX](#). It is made up of loosely coupled libraries, which are showcased with end-to-end integrated [guides](#) and [examples](#).

## jax.jit

```
jax.jit(fun, in_shardings=UnspecifiedValue, out_shardings=UnspecifiedValue,  
static_argnums=None, static_argnames=None, donate_argnums=None, donate_argnames=None,  
keep_unused=False, device=None, backend=None, inline=False, abstracted_axes=None)
```

Sets up `fun` for just-in-time compilation with XLA.

[\[source\]](#)


## GPU memory allocation

JAX will **preallocate 75% of the total GPU memory when the first JAX operation is run**. Preallocating minimizes allocation overhead and memory fragmentation, but can sometimes cause out-of-memory (OOM) errors. If your JAX process fails with OOM, the following environment variables can be used to override the default behavior:


```
XLA_PYTHON_CLIENT_PREALLOCATE=false
```


This disables the preallocation behavior. JAX will instead allocate GPU memory as needed, potentially decreasing the overall memory usage. However, this behavior is more prone to GPU memory fragmentation, meaning a JAX program that uses most of the available GPU memory may OOM with preallocation disabled.

## On CPU



## On GPU





Node 8 × A100 80Go

Really worth it ?

