



# Deep Learning Optimisé - Jean Zay

---

Best Practices and State Of The Art



INSTITUT DU  
DÉVELOPPEMENT ET DES  
RESSOURCES EN  
INFORMATIQUE  
SCIENTIFIQUE



# Fast.ai tips and engineering

fast.ai

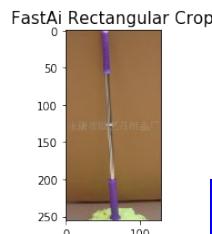
“An AI speed test shows clever coders can still beat tech giants like Google and Intel.” DAWN Bench competition 2018



OneCycle lr scheduler  
+ lr finder



Popularizes the works of  
[Leslie N. Smith](#)

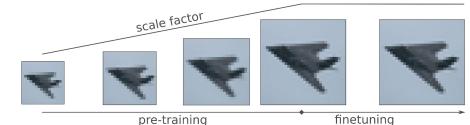


Thanks to Global Average Pooling



Test Rectangular  
Validation Technique

Progressive image  
resizing



Dynamic batch size



# ML Perf - Reference for AI Supercomputing



Fair and useful benchmarks for measuring training and inference performance of ML hardware, software, and services.

Industry standard

- Hardware
- Framework
- SOTA

## Training



Speech Recognition  
RNN-T



NLP  
BERT



Recommender  
DLRM



Reinforcement Learning  
MiniGo



Biomedical Image Segmentation  
UNet-3D



Object Detection (Light weight)  
SSD



Object Detection (Heavy weight)  
Mask R-CNN



Image Classification  
ResNet-50 v1.5

## Industry-Standard Generative AI Training Benchmarks

MLPerf Training v3.1



GPT-3 175B  
Large Language Model



Stable Diffusion  
Text-to-Image



DLRMv2  
Recommendation



BERT-Large  
NLP



RetinaNet  
Object Detection,  
Lightweight



Mask R-CNN  
Object Detection,  
Heavyweight



3D U-Net  
Biomedical Image  
Segmentation



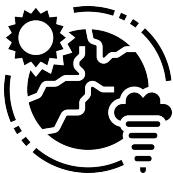
RNN-T  
Speech Recognition



ResNet-50 v1.5  
Image Classification



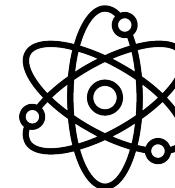
## Training HPC



Climate segmentation  
**DeepCAM**



Cosmology parameter prediction  
**CosmoFlow**



Quantum molecular modeling  
**DimNet++**

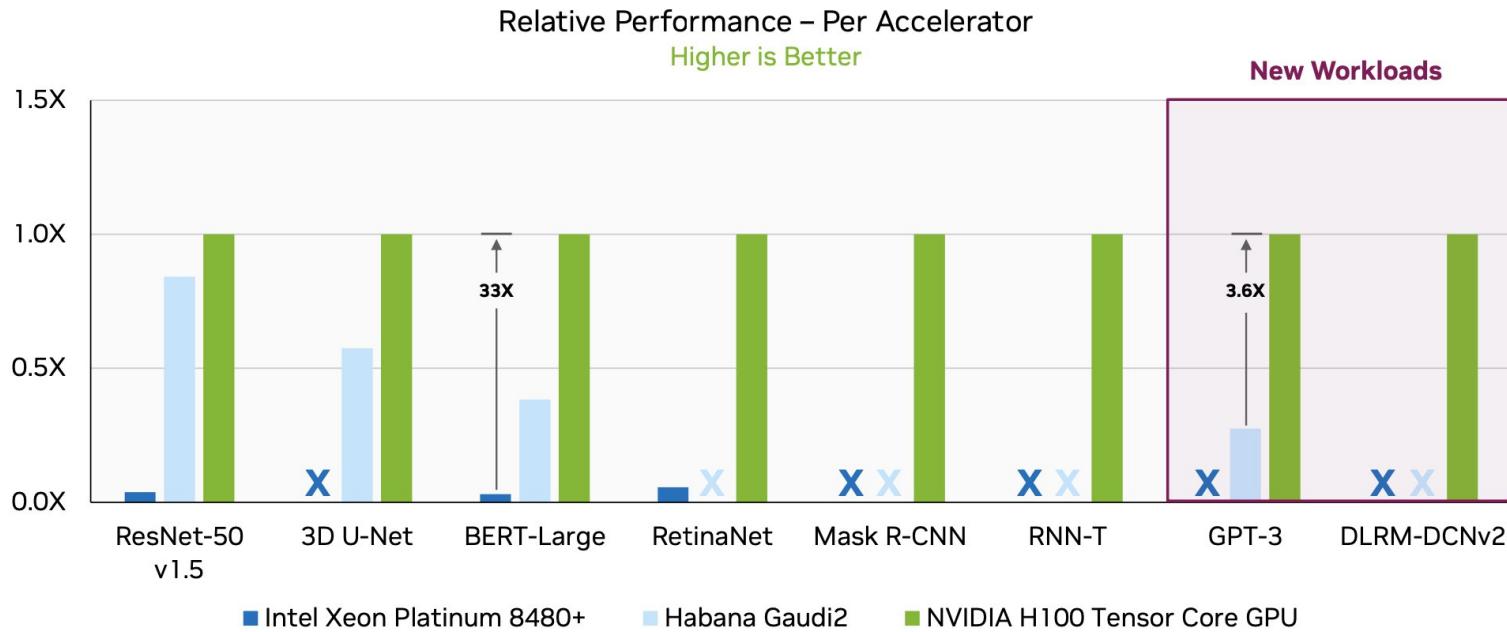
## Inference :



- Datacenter
- Edge
- Mobile
- Tiny

## NVIDIA H100 GPU Extends AI Training Leadership

Fastest and most versatile AI accelerator

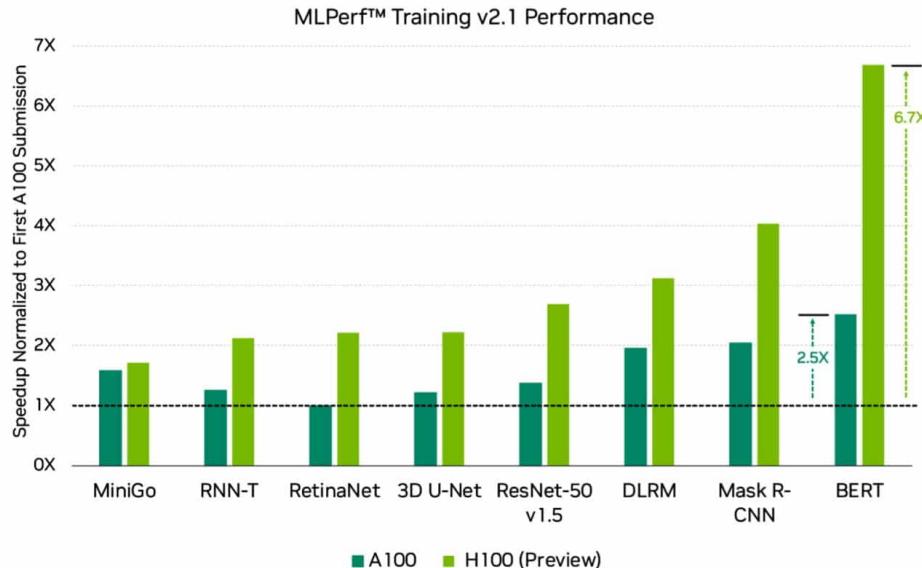


# ML Perf - Evolution



## NVIDIA AI and H100 Deliver 6.7X in 2.5 Years

Full-stack innovation fuels continuous performance gains



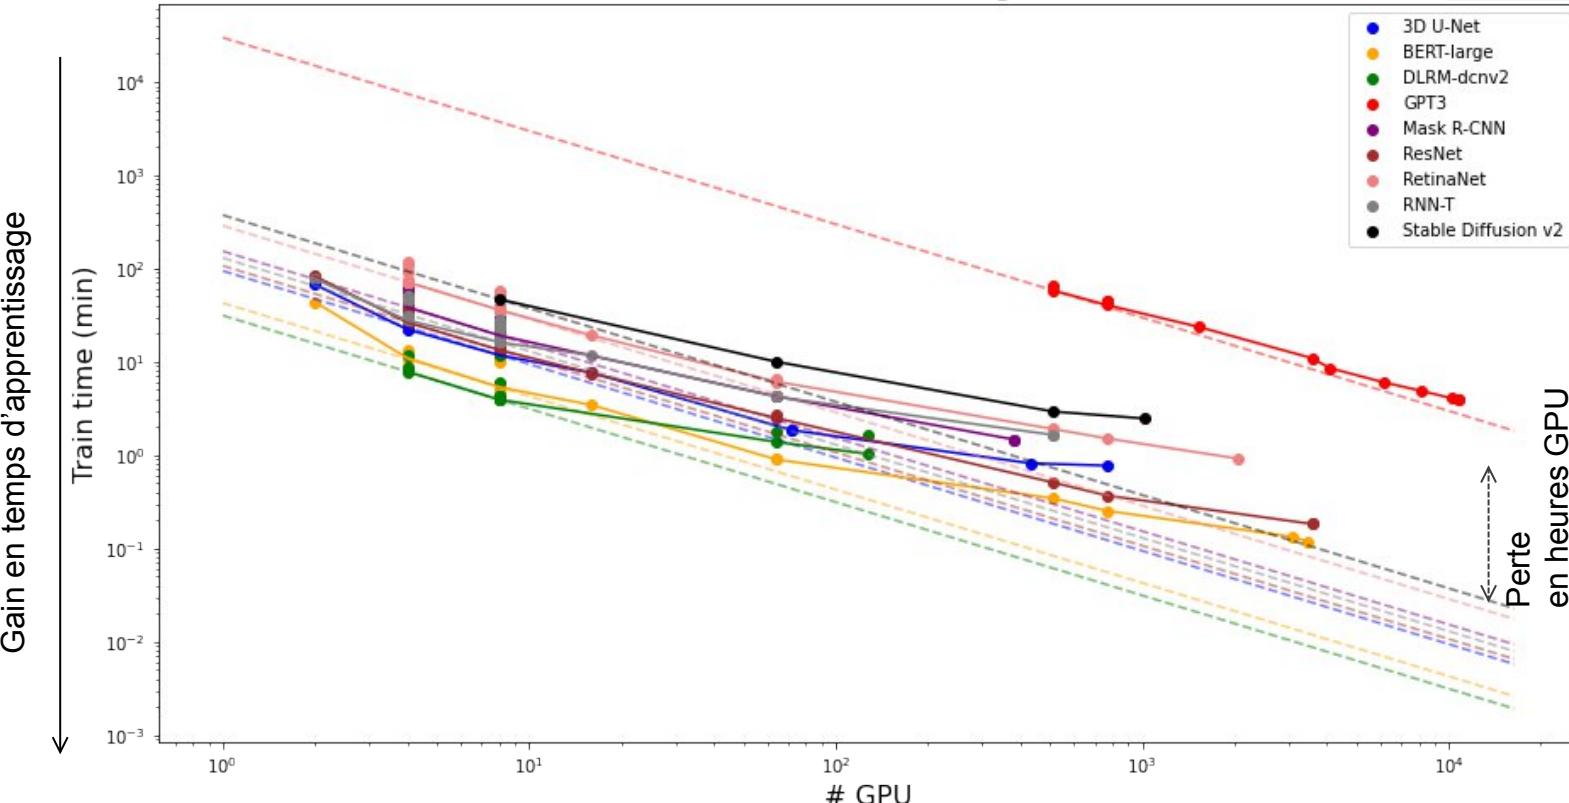
Up to  
**6.7X**  
Higher performance with new H100 GPUs

Up to  
**2.5X**  
Speedup on existing A100 GPUs with software

ResNet-50 v1.5: 8x NVIDIA 0.7-18, 8x NVIDIA 2.1-2060, 8x NVIDIA 2.1-2091 | BERT: 8x NVIDIA 0.7-19, 8x NVIDIA 2.1-2062, 8x NVIDIA 2.1-2091 | DLRM: 8x NVIDIA 0.7-17, 8x NVIDIA 2.1-2059, 8x NVIDIA 2.1-2091 | Mask R-CNN: 8x NVIDIA 0.7-19, 8x NVIDIA 2.1-2062, 8x NVIDIA 2.1-2091 | RetinaNet: 8x NVIDIA 2.0-2091, 8x NVIDIA 2.1-2061, 8x NVIDIA 2.1-2091 | RNN-T: 8x NVIDIA 1.0-1060, 8x NVIDIA 2.1-2061, 8x NVIDIA 2.1-2091 | Mini Go: 8x NVIDIA 0.7-20, 8x NVIDIA 2.1-2063, 8x NVIDIA 2.1-2091 | 3D U-Net: 8x NVIDIA 1.0-1059, 8x NVIDIA 2.1-2060, 8x NVIDIA 2.1-2091

First NVIDIA A100 Tensor Core GPU results normalized for throughput due to higher accuracy requirements introduced in MLPerf™ Training 2.0 where applicable.  
MLPerf™ name and logo are trademarks. See [www.mlperf.org](http://www.mlperf.org) for more information.

MLPerf - H100 - Training v3.0



# AlgoPerf: Training Algorithms benchmark

Speeding up training requires new competitive benchmarks to detect robust improvement :

Hyperparameter	AdamW	NadamW	Heavy Ball	Nesterov
Base LR	Log [1e-4,1e-2]	Log [1e-4,1e-2]	Log [1e-1,10]	Log [1e-1,10]
Weight decay	Log [5e-3,1]	Log [5e-3,1]	Log [1e-7,1e-5]	Log [1e-7,1e-5]
1 - $\beta_1$ [Default]	0.1	0.1	0.1	0.1
1 - $\beta_1$ [Tuned]	Log [2e-2,0.5]	Log [4e-3,0.1]	Log [5e-3,0.3]	Log [5e-3,0.3]
1 - $\beta_2$	0.999	0.999	NA	NA
Schedule	warmup + cosine decay	warmup + cosine decay	warmup + linear decay	warmup + linear decay
Warmup	5%	5%	5%	5%
Decay factor	NA	NA	{1e-2,1e-3}	{1e-2,1e-3}
Decay steps	NA	NA	0.9	0.9
Label smoothing	{0.1,0.2}	{0.1,0.2}	{0.1,0.2}	{0.1,0.2}
Dropout (Tied)	{0.0,0.1}	{0.0,0.1}	{0.0,0.1}	{0.0,0.1}

Hyperparameter	LAMB	Adafactor	SAM(w. Adam)	Distributed Shampoo
Base LR	Log [1e-4,1e-2]	Log [1e-4,1e-2]	Log [1e-4,1e-2]	Log [1e-4,1e-2]
Weight decay	Log [1e-3,1]	Log [1e-3,1]	Log [1e-2,0.2]	Log [5e-3,1]
1 - $\beta_1$ [Default]	0.1	0.1	0.1	0.1
1 - $\beta_1$ [Tuned]	Log [2e-2,0.5]	Log [1e-2,0.45]	Log [5e-2,0.43]	Log [1e-2,0.15]
1 - $\beta_2$	0.999	0.999	0.999	0.999
$\rho$	NA	NA	{0.01, 0.02, 0.05}	NA
Schedule	warmup + cosine decay	warmup + cosine decay	warmup + cosine decay	warmup + cosine decay
Warmup	5%	5%	5%	5%
Decay factor	NA	NA	NA	NA
Decay steps	NA	NA	NA	NA
Label smoothing	{0.1,0.2}	{0.1,0.2}	{0.1,0.2}	{0.1,0.2}
Dropout (Tied)	{0.0,0.1}	{0.0,0.1}	{0.0,0.1}	{0.0,0.1}

- **Activation function:** Most base workloads employed ReLU as the activation function and we explored alternative activation functions such as GELU, SiLU, or TanH.

- **Pre-LN vs Post-LN:** For Transformer-based models, the base workload was usually the PRE-LAYER NORM (PRE-LN) (Xiong et al., 2020) version. We changed these to POST-LAYER NORM (POST-LN, see Figure 3).

- **Attention temperature:** For the WMT TRANSFORMER, we modified the attention layers to compute Softmax  $\left( \frac{c X W^Q (X W^K)^T}{\sqrt{D/H}} \right)$  where  $c$  is a constant scalar denoting the attention temperature. The default self-attention implementation sets  $c = 1$ . In order to artificially induce instabilities similar to those faced by larger versions of these models, we set  $c = 1.6$ .

**Initialization scales:** For the DLRM model, changing the scale of the initial weights of the embedding layer resulted in a variant. For the RESNET model, changing the initial batch normalization layer scale weights resulted in a workload variant.

**Normalization layer:** We changed the type of normalization layer employed in the model. Common changes included interchanging batch normalization with layer normalization, as well as instance normalization with layer normalization.

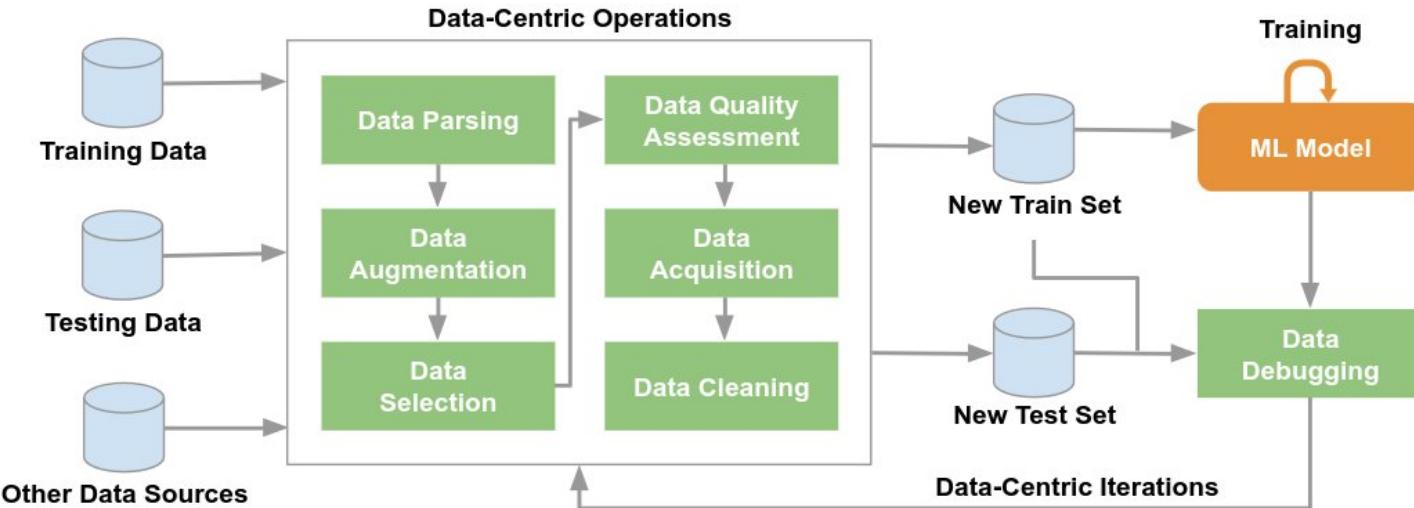
**Width, depth, and channels:** We explored changing model width, depth, and number of channels, as applicable.

**Input pipeline:** On LIBRISPEECH, we found changing SPECAUGMENT strength to be an effective strategy.

**Residual connection structure and scaling:** For the DLRM model, we created a variant with additional residual connections. For DEEPSPEECH, we removed residual connections from the model.

**Pooling layer type:** Changing the pooling layer type from *global average pooling* to *max average pooling* resulted in a variant for the ViT workload.

# DataPerf – Data-Centric benchmark



**Figure 1:** Typical benchmarks are model-centric, and therefore focus on the model design and training stages of the ML pipeline (shown in orange). However, to develop high-quality ML applications, users often employ a collection of data-centric operations to improve data quality and repeated data-centric iterations to refine these operations. DataPerf aims to benchmark all major stages of such a data-centric pipeline (shown in green) to improve ML data quality.

# Performance Tuning Guide

- Enable asynchronous data loading and augmentation

```
torch.utils.data.DataLoader  
    num_workers > 0  
    pin_memory=True
```

- Disable gradient calculation for validation or inference

```
with torch.no_grad():  
    val_outputs = model(val_images)  
    val_loss = criterion(val_outputs, val_labels)
```

- Use mixed precision and AMP

```
from torch.cuda.amp import autocast, GradScaler  
with autocast():
```

- Use efficient data-parallel backend

```
torch.nn.parallel.DistributedDataParallel
```

- Disable bias for convolutions directly followed by a batch norm

```
nn.Conv2d(..., bias=False, ....)
```

Models available from `torchvision` already implement this optimization.

- Enable `channels_last` memory format for computer vision models

```
x = x.to(memory_format=torch.channels_last)
```

- Disable debugging APIs

```
anomaly detection: torch.autograd.detect_anomaly or torch.autograd.set_detect_anomaly(True)  
profiler related: torch.autograd.profiler.emit_nvtx, torch.autograd.profiler.profile  
autograd gradcheck: torch.autograd.gradcheck or torch.autograd.gradgradcheck
```

- Create tensors directly on the target device

```
torch.rand(size).cuda()
```

```
torch.rand(size, device='cuda')
```

- Fuse pointwise operations

Pointwise operations (elementwise addition, multiplication, math functions - sin(), cos(), sigmoid() etc.) can be fused into a single kernel to amortize memory access time and kernel launch time. **PyTorch JIT** can fuse kernels automatically.

```
@torch.jit.script
def fused_gelu(x):
    return x * 0.5 * (1.0 + torch.erf(x / 1.41421))
```

- Enable cuDNN auto-tuner

For convolutional networks

```
torch.backends.cudnn.benchmark = True
```

- Avoid unnecessary CPU-GPU synchronization

```
print(cuda_tensor)
cuda_tensor.item()
memory copies: tensor.cuda(), cuda_tensor.cpu() and equivalent tensor.to(device) calls
cuda_tensor.nonzero()
python control flow e.g. if (cuda_tensor != 0).all()
```

- Load-balance workload in a distributed setting

The core idea is to **distribute workload over all workers** as uniformly as possible within **each global batch**. For example Transformer solves imbalance by **forming batches with approximately constant number of tokens** (and variable number of sequences in a batch), other models solve imbalance by **bucketing samples with similar sequence length** or even by **sorting dataset** by sequence length.

- Preallocate memory in case of variable input length

For Speech Recognition or NLP, **preeexecute** a forward and a backward pass with a **generated batch of inputs with maximum sequence length** (either corresponding to max length in the training dataset or to some predefined threshold). This step **preallocates buffers** of maximum size, which can be reused in subsequent training iterations.

- Match the order of layers in constructors and during the execution if using `DistributedDataParallel``(find_unused_parameters=True)`

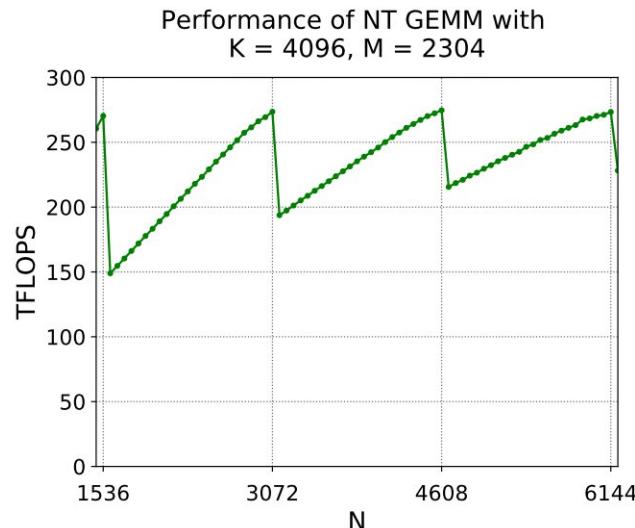
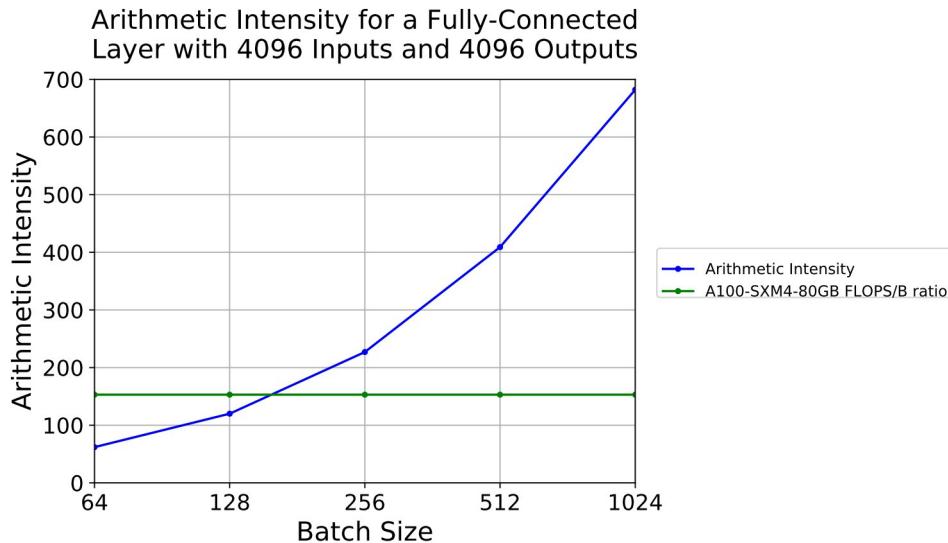
To maximize the amount of overlap, the **order in model constructors** should roughly **match** the order during the execution. If the order doesn't match, then **all-reduce** for the entire bucket **waits** for the gradient which is the last to arrive.

With `find_unused_parameters=False` it's **not necessary** to reorder layers or parameters to achieve optimal performance.

# Deep Learning Performance Documentation



$$\text{Arithmetic Intensity} = \frac{\text{number of FLOPS}}{\text{number of byte accesses}} = \frac{2 \cdot (M \cdot N \cdot K)}{2 \cdot (M \cdot K + N \cdot K + M \cdot N)} = \frac{M \cdot N \cdot K}{M \cdot K + N \cdot K + M \cdot N}$$



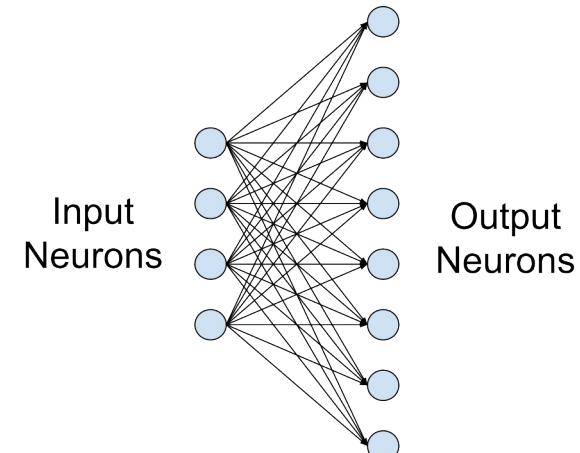
Wave Quantization effect

# Linear/Fully-Connected Layers User's Guide



The following quick start checklist provides specific tips for **fully-connected layers**.

- Choose the batch size and the number of inputs and outputs to be **divisible by 4 (TF32) / 8 (FP16) / 16 (INT8)** to run efficiently on **Tensor Cores**. For best efficiency on **A100**, choose these parameters to be **divisible by 32 (TF32) / 64 (FP16) / 128 (INT8)** .
- Especially when ones are small, choosing the batch size and the number of inputs and outputs to be **divisible by at least 64** and **ideally 256** can streamline tiling and reduce overhead.
- **Larger values** for batch size and the number of inputs and outputs **improve** parallelization and efficiency.
- As a rough guideline, choose batch sizes and neuron counts **greater than 128** to avoid being limited by memory bandwidth.

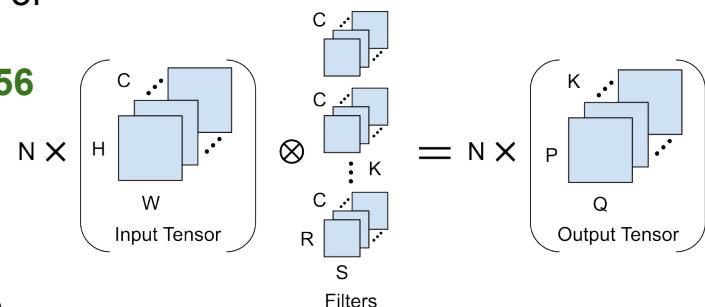


# Convolutional Layers User's Guide



The following quick start checklist provides specific tips for **convolutional layers**.

- Choose the number of **input and output channels** to be divisible by **8 (for FP16) or 4 (for TF32)** to run efficiently on **Tensor Cores**. For the **first convolutional layer** in most CNNs with **3-channel** images, **padding to 4 channels** is sufficient if a stride of 2 is used.
- Choose parameters to be divisible by **at least 64** and **ideally 256** to enable efficient tiling and reduce overhead.
- **Larger values** for size-related parameters can improve parallelization.
- When the **size of the input is the same** in each iteration, **autotuning** is an efficient method to ensure the selection of the ideal algorithm for each convolution in the network.  
`torch.backends.cudnn.benchmark = True.`
- Choose tensor layouts in memory to avoid transposing input and output data. We recommend using the **NHWC format** where possible.

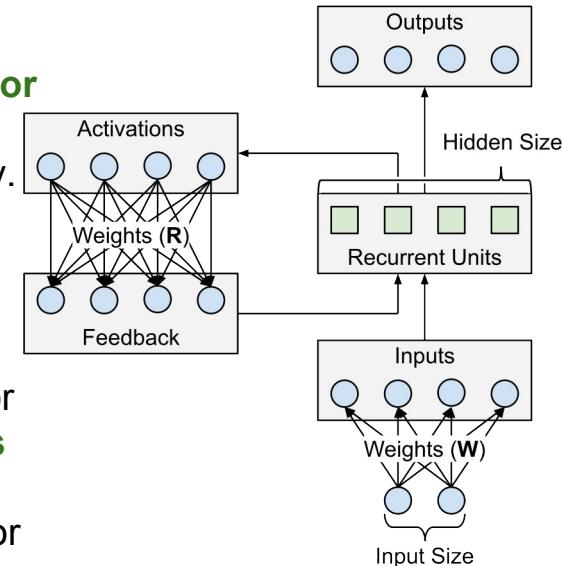


# Recurrent Layers User's Guide



The following quick start checklist provides specific tips for **recurrent layers**.

- **Recurrent operations can be parallelized.** We recommend using NVIDIA® cuDNN implementations, which do this automatically.
- When using the **standard implementation**, minibatch size and hidden sizes should be:
  - **Divisible by 8 (for FP16) or 4 (for TF32)** to run efficiently on **Tensor Cores**.
  - **Divisible by at least 64 and ideally 256** to improve tiling efficiency.
  - **Greater than 128 (minibatch size) or 256 (hidden sizes)** to be limited by computation rate rather than memory bandwidth.
- When using the **persistent implementation** (available for FP16 data only):
  - **Hidden sizes** should be **divisible by 32** to run efficiently on Tensor Cores. Better tiling efficiency may be achieved **by larger multiples of 2, up to 256**.
  - **Minibatch size** should be **divisible by 8** to run efficiently on Tensor Cores...
  - **Try increasing parameters for better efficiency.**



# Memory-Limited Layers User's Guide



The following quick start checklist provides specific tips **for layers whose performance is limited by memory accesses (Batch Normalization, Activations, Pooling, ...)**.

- Explore the available implementations of each layer in the **NVIDIA cuDNN API Reference** or your framework. Often the best way to improve performance is to choose **a more efficient implementation**.
- **Be aware of the number of memory accesses** required for each layer. Performance of a memory-bound calculation is simply based on the number of inputs, outputs, and weights that need to be loaded and/or stored per pass. We don't have recommended parameter tweaks for these layers.
- **Be aware of the impact of each layer** on the overall training step performance. **Memory-bound layers** are most likely to take a significant amount of time in small networks where there are no large and computation-heavy layers to dominate performance.

# Hugging Face

Hugging Face Search models, datasets, users... Models Datasets Spaces Posts Docs Solutions Pricing

Transformers v4.4.1.3 EN 127,323 Interoperability with GGUF files

Performance and Scalability Overview LLM inference optimization Quantization

Efficient Training Techniques Methods and tools for efficient training on a single GPU Multiple GPUs and parallelism Fully Sharded Data Parallel DeepSpeed Efficient training on CPU Distributed CPU training Training on TPU with TensorFlow PyTorch training on Apple silicon Custom hardware for training Hyperparameter Search using Trainer API

Optimizing Inference CPU inference GPU inference Instantiate a big model

## Efficient Training on Multiple GPUs

If training a model on a single GPU is too slow or if the model's weights do not fit in a single GPU's memory, transitioning to a multi-GPU setup may be a viable option. Prior to making this transition, thoroughly explore all the strategies covered in the [Methods and tools for efficient training on a single GPU](#) as they are universally applicable to model training on any number of GPUs. Once you have employed those strategies and found them insufficient for your case on a single GPU, consider moving to multiple GPUs.

Transitioning from a single GPU to multiple GPUs requires the introduction of some form of parallelism, as the workload must be distributed across the resources. Multiple techniques can be employed to achieve parallelism, such as data parallelism, tensor parallelism, and pipeline parallelism. It's important to note that there isn't a one-size-fits-all solution, and the optimal settings depend on the specific hardware configuration you are using.

This guide offers an in-depth overview of individual types of parallelism, as well as guidance on ways to combine techniques and choosing an appropriate approach. For step-by-step tutorials on distributed training, please refer to the [Accelerate documentation](#).

While the main concepts discussed in this guide are likely applicable across frameworks, here we focus on PyTorch-based implementations.

Before diving deeper into the specifics of each technique, let's go over the rough decision process when training large models on a large infrastructure.

### Scalability strategy

Begin by estimating how much vRAM is required to train your model. For models hosted on the 😊 Hub, use our [Model Memory](#)

Efficient Training on Multiple GPUs

Scalability strategy

Data Parallelism

DataParallel vs DistributedData

Parallel

ZeRO Data Parallelism

From Naive Model Parallelism to Pipeline Parallelism

Tensor Parallelism

Data Parallelism + Pipeline Parallelism

Data Parallelism + Pipeline Parallelism + Tensor Parallelism

ZeRO Data Parallelism + Pipeline Parallelism + Tensor Parallelism

FlexFlow

GPU selection

Number of GPUs

Order of GPUs